

Retrieval-Augmented Generation (RAG) Integration

Overview

This document outlines the strategy for implementing Retrieval-Augmented Generation (RAG) in the Platform Dashboard. RAG enhances the AI chatbot capabilities by retrieving relevant information from a knowledge base before generating responses, improving accuracy and relevance while reducing hallucinations.

Table of Contents

- 1. [Introduction](#)
- 2. [Architecture](#)
- 3. [Database Schema Updates](#)
- 4. [Implementation Steps](#)
- 5. [Vector Database Integration](#)
- 6. [Document Processing Pipeline](#)
- 7. [Frontend Integration](#)
- 8. [Backend Services](#)
- 9. [API Endpoints](#)
- 10. [Testing and Validation](#)
- 11. [Future Enhancements](#)

Introduction

Retrieval-Augmented Generation (RAG) combines retrieval-based and generation-based approaches for AI responses:


- 1. **Retrieval:** When a user submits a query, the system retrieves relevant information from a knowledge base
- 2. **Augmentation:** This retrieved information is combined with the user's query
- 3. **Generation:** The augmented context is sent to the LLM for response generation

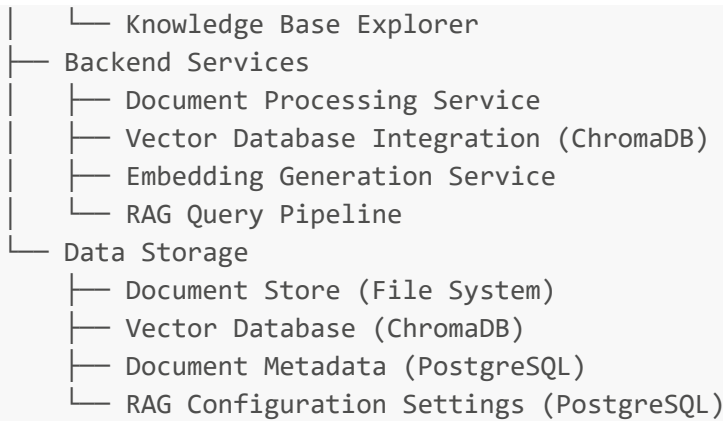
Benefits:

- Enhanced factual accuracy using up-to-date information
- Reduced hallucinations by grounding responses in sourced content
- Ability to cite sources of information
- Improved handling of domain-specific queries

Architecture

The RAG implementation will be built on our existing Ollama integration, with the following components:





Database Schema Updates

New Tables

1. **vector_stores**

- Track configured vector database instances

```
CREATE TABLE public.vector_stores (
    id UUID DEFAULT gen_random_uuid() NOT NULL PRIMARY KEY,
    name VARCHAR(255) NOT NULL,
    store_type VARCHAR(50) NOT NULL,
    connection_string TEXT,
    is_active BOOLEAN DEFAULT TRUE,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    config JSONB
);
```

2. **document_collections**

- Organize documents into collections (e.g., "Technical Documentation", "Policies")

```
CREATE TABLE public.document_collections (
    id UUID DEFAULT gen_random_uuid() NOT NULL PRIMARY KEY,
    name VARCHAR(255) NOT NULL,
    description TEXT,
    user_id UUID NOT NULL REFERENCES public.users(id) ON DELETE CASCADE,
    vector_store_id UUID REFERENCES public.vector_stores(id) ON DELETE SET
    NULL,
    embedding_model VARCHAR(255),
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    is_active BOOLEAN DEFAULT TRUE,
    metadata JSONB
);
```

3. documents

- Enhanced version of existing `pdfs` table to support all document types

```
CREATE TABLE public.documents (  
  id UUID DEFAULT gen_random_uuid() NOT NULL PRIMARY KEY,  
  user_id UUID NOT NULL REFERENCES public.users(id) ON DELETE CASCADE,  
  collection_id UUID REFERENCES public.document_collections(id) ON DELETE  
SET NULL,  
  title VARCHAR(255) NOT NULL,  
  file_path TEXT NOT NULL,  
  file_name VARCHAR(255) NOT NULL,  
  file_type VARCHAR(50) NOT NULL,  
  content_text TEXT,  
  content_hash VARCHAR(64),  
  vector_id VARCHAR(255),  
  processing_status VARCHAR(20) DEFAULT 'pending' NOT NULL, -- 'pending',  
  'processing', 'processed', 'failed'  
  is_indexed BOOLEAN DEFAULT FALSE,  
  chunk_count INTEGER DEFAULT 0,  
  uploaded_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,  
  processed_at TIMESTAMP,  
  indexed_at TIMESTAMP,  
  last_accessed_at TIMESTAMP,  
  metadata JSONB  
);
```

4. document_chunks

- Store individual chunks of documents for fine-grained retrieval

```
CREATE TABLE public.document_chunks (  
  id UUID DEFAULT gen_random_uuid() NOT NULL PRIMARY KEY,  
  document_id UUID NOT NULL REFERENCES public.documents(id) ON DELETE  
CASCADE,  
  chunk_index INTEGER NOT NULL,  
  content TEXT NOT NULL,  
  vector_id VARCHAR(255),  
  embedding VECTOR(1536),  
  token_count INTEGER,  
  created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,  
  metadata JSONB  
);
```

5. rag_settings

- Configure RAG behavior and parameters

```
CREATE TABLE public.rag_settings (  
  id SERIAL PRIMARY KEY,  
  embedding_model VARCHAR(255) DEFAULT 'ollama/nomic-embed-text',  
  chunk_size INTEGER DEFAULT 1000,  
  chunk_overlap INTEGER DEFAULT 200,  
  similarity_top_k INTEGER DEFAULT 4,  
  search_type VARCHAR(20) DEFAULT 'similarity',  
  created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,  
  updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,  
  config JSONB  
);
```

Updating Existing Tables

1. **messages** - Add references to retrieved document chunks

```
ALTER TABLE public.messages ADD COLUMN retrieved_chunks JSONB;
```

2. **chat_sessions** - Add option to enable/disable RAG for a session

```
ALTER TABLE public.chat_sessions ADD COLUMN use_rag BOOLEAN DEFAULT TRUE;  
ALTER TABLE public.chat_sessions ADD COLUMN rag_collections JSONB;
```

Implementation Steps

Phase 1: Foundation (Weeks 1-2)

1. Database Setup

- Create new tables in PostgreSQL
- Create migration scripts
- Update database documentation

2. Vector Database Integration

- Install ChromaDB locally or via Docker
- Implement connection service
- Create CRUD operations for vectors
- Implement configuration management

3. Document Processing Service

- Create file upload and processing pipeline
- Implement text extraction from PDFs, DOCXs, etc.
- Implement text chunking strategies
- Create embedding generation service using Ollama models

Phase 2: Core RAG Implementation (Weeks 3-4)

1. Query Pipeline

- Implement query embedding generation
- Develop vector search functionality
- Create context assembly for LLM prompts
- Implement citation tracking

2. Ollama Integration

- Enhance Ollama service to support RAG
- Implement prompt templates for RAG queries
- Create specialized chat completion with context

3. API Endpoints

- Implement document management endpoints
- Create collection management endpoints
- Develop RAG-enabled chat endpoints

Phase 3: Frontend and User Experience (Weeks 5-6)

1. Document Management UI

- Create document upload interface
- Implement collection management
- Develop processing status indicators

2. Enhanced Chat Interface

- Update chat UI to show citations
- Implement RAG toggle option
- Create source preview functionality

3. Knowledge Base Explorer

- Develop interface to browse uploaded documents
- Create collection and document search
- Implement document preview

Vector Database Integration

ChromaDB will be used as the vector database for storing and retrieving document embeddings:

1. Installation

```
npm install chromadb
```

2. Configuration

- Local ChromaDB instance (development)
- Persistent store using PersistentClient
- Collection per document group

3. Integration Points

- Document upload pipeline (embedding generation)
- Query processing (similarity search)
- Collection management (CRUD operations)

Core Functionality

```
// src/services/vectorStoreService.js
const { ChromaClient } = require('chromadb');

class VectorStoreService {
  constructor(config) {
    this.client = new ChromaClient(config.host, config.port);
    this.defaultEmbeddingFunction =
this.getEmbeddingFunction(config.embeddingModel);
  }

  async getCollection(name, embeddingFunction = this.defaultEmbeddingFunction) {
    try {
      return await this.client.getCollection(name, embeddingFunction);
    } catch (error) {
      return await this.client.createCollection(name, { embeddingFunction });
    }
  }

  getEmbeddingFunction(modelName) {
    // Integration with Ollama embedding models
    return {
      generate: async (texts) => {
        // Call Ollama embedding endpoint
        const embeddings = await ollamaService.generateEmbeddings(modelName,
texts);
        return embeddings;
      }
    };
  }

  // Other methods for CRUD operations, search, etc.
}
```

Document Processing Pipeline

The document processing pipeline will handle various document types and prepare them for RAG:

Directory Structure

```

productdemo/
├── Documents/           # Storage for uploaded document files
│   ├── user_id_1/      # Organized by user ID
│   │   ├── doc_id_1/   # Each document gets its own folder
│   │   │   ├── original.pdf
│   │   │   └── doc_id_2/
│   │   │       └── original.docx
│   │   └── user_id_2/
│   │       └── ...
│   └── ...
├── Embeddings/         # Storage for document embeddings
│   ├── user_id_1/      # Organized by user ID
│   │   ├── doc_id_1/   # Each document gets its own folder
│   │   │   ├── metadata.json # Document metadata including hash
│   │   │   ├── chunks.json  # Text chunks
│   │   │   └── embeddings.bin # Binary embedding vectors
│   │   └── doc_id_2/
│   │       └── ...
│   └── user_id_2/
│       └── ...
└── ...

```

Processing Flow

1. Upload Phase

- File validation and sanitization
- Document metadata extraction
- Storage in file system (`Documents/user_id/doc_id/original.[ext]`)
- Entry created in database with status "pending"

2. Processing Phase

- Text extraction based on document type
- Content chunking with configurable overlap
- Embedding generation using Ollama models
- Vector storage in ChromaDB
- Status updates in database
- Storage of embeddings in `Embeddings/user_id/doc_id/` directory

3. Indexing Phase

- Metadata indexing
- Vector indexing optimization
- Relationship mapping for collections

Deduplication Strategy

To avoid reprocessing identical documents:

1. Hash Verification

- When a document is uploaded, calculate its content hash
- Check if a document with the same hash exists in the user's collection
- If match found, verify metadata similarity (optional)

2. Reuse Existing Embeddings

- If duplicate detected, link to existing embeddings in database
- Update document record with reference to existing vector IDs
- Skip processing pipeline, mark as "processed" immediately

3. Partial Updates

- For similar but not identical documents, consider partial updates
- Only process and embed changed sections (future enhancement)

Text Extraction and Chunking

```
// src/services/documentProcessingService.js
const fs = require('fs');
const path = require('path');
const pdf = require('pdf-parse');
const mammoth = require('mammoth');

class DocumentProcessingService {
  constructor(config, vectorStoreService) {
    this.config = config;
    this.vectorStoreService = vectorStoreService;
    this.chunkSize = config.chunkSize || 1000;
    this.chunkOverlap = config.chunkOverlap || 200;
  }

  async processDocument(document) {
    // Extract text based on document type
    const text = await this.extractText(document);

    // Chunk text into segments
    const chunks = this.chunkText(text);

    // Generate embeddings and store in vector db
    await this.vectorizeChunks(document.id, chunks, document.collection_id);

    return {
      documentId: document.id,
      chunkCount: chunks.length,
      status: 'processed'
    };
  }

  async extractText(document) {
    const filePath = document.file_path;
    const fileType = document.file_type.toLowerCase();
```



```
switch (fileType) {
  case 'pdf':
    return await this.extractFromPDF(filePath);
  case 'docx':
    return await this.extractFromDOCX(filePath);
  case 'txt':
    return await this.extractFromTXT(filePath);
  // Add other document types as needed
  default:
    throw new Error(`Unsupported file type: ${fileType}`);
}
}

// Methods for specific file type extraction
async extractFromPDF(filePath) {
  const dataBuffer = fs.readFileSync(filePath);
  const data = await pdf(dataBuffer);
  return data.text;
}

async extractFromDOCX(filePath) {
  const result = await mammoth.extractRawText({ path: filePath });
  return result.value;
}

async extractFromTXT(filePath) {
  return fs.readFileSync(filePath, 'utf8');
}

// Methods for chunking
chunkText(text) {
  const chunks = [];
  let startIndex = 0;

  while (startIndex < text.length) {
    // Calculate end index with consideration for overlap
    const endIndex = Math.min(startIndex + this.chunkSize, text.length);
    chunks.push(text.slice(startIndex, endIndex));

    // Move start index forward, accounting for overlap
    startIndex += this.chunkSize - this.chunkOverlap;

    // Ensure we make progress even with large overlap
    if (startIndex <= (endIndex - this.chunkSize/2)) {
      startIndex = endIndex - this.chunkOverlap;
    }
  }

  return chunks;
}

// Methods for vectorization and deduplication
async vectorizeChunks(documentId, chunks, collectionId) {
  // Generate embeddings for each chunk
```

```
    const embeddings = await this.generateEmbeddings(chunks);

    // Store chunks and embeddings
    await this.storeChunksAndEmbeddings(documentId, chunks, embeddings);

    // Update document status
    await this.updateDocumentStatus(documentId, chunks.length);

    return chunks.length;
}

async checkForDuplicates(document) {
    // Calculate content hash if not already done
    if (!document.content_hash) {
        const text = await this.extractText(document);
        document.content_hash = this.calculateContentHash(text);
        await this.updateDocumentHash(document.id, document.content_hash);
    }

    // Check for existing document with same hash
    const existingDoc = await this.findDocumentByHash(
        document.content_hash,
        document.user_id
    );

    if (existingDoc) {
        console.log(`Duplicate document detected: ${document.id} matches ${existingDoc.id}`);

        // Link to existing embeddings
        await this.linkToExistingEmbeddings(document.id, existingDoc.id);

        // Mark as processed without running the pipeline
        await this.updateDocumentStatus(
            document.id,
            existingDoc.chunk_count,
            'processed',
            'Used existing embeddings from duplicate document'
        );

        return true; // Duplicate found
    }

    return false; // No duplicate found
}

calculateContentHash(text) {
    const crypto = require('crypto');
    return crypto.createHash('sha256').update(text).digest('hex');
}
```

Frontend Integration

The frontend will be updated to support RAG features:

1. Document Management Page

- Upload interface with drag-and-drop
- Collection management
- Processing status and statistics

2. Chat Enhancements

- Toggle for RAG mode
- Collection selector for context scope
- Citation display in messages
- Source preview on citation click

3. Knowledge Base Explorer

- Document browser with search and filters
- Collection management interface
- Document preview and metadata display

Component Updates

```
// client/src/components/chat/ChatInput.tsx
// Add RAG options to ChatInput component

const ChatInput: React.FC<ChatInputProps> = ({
  onSendMessage,
  isLoading,
  useRAG,
  onToggleRAG,
  selectedCollections,
  onSelectCollections
}) => {
  // Existing code...

  return (
    <div className="chat-input-container">
      {/* Existing input elements */}

      <div className="rag-controls">
        <Switch
          isChecked={useRAG}
          onChange={onToggleRAG}
          label="Use Knowledge Base"
        />

        {useRAG && (
          <CollectionSelector
            selectedCollections={selectedCollections}
          />
        )}
      </div>
    </div>
  )
}
```

```
        onChange={onSelectCollections}
      />
    )}
  </div>
</div>
);
};
```

Backend Services

The backend will require several new services:

1. **ragService.js**: Orchestrates the RAG process
2. **vectorStoreService.js**: Manages vector database operations
3. **documentProcessingService.js**: Handles document processing
4. **embeddingService.js**: Generates embeddings using Ollama

RAG Service

```
// src/services/ragService.js
class RAGService {
  constructor(
    vectorStoreService,
    documentService,
    ollamaService,
    config
  ) {
    this.vectorStoreService = vectorStoreService;
    this.documentService = documentService;
    this.ollamaService = ollamaService;
    this.config = config;
  }

  async queryWithRAG(query, model, collectionIds, options = {}) {
    // Generate embedding for the query
    const queryEmbedding = await this.getQueryEmbedding(query);

    // Retrieve relevant chunks from vector store
    const relevantChunks = await this.retrieveRelevantChunks(
      queryEmbedding,
      collectionIds,
      options.topK || 4
    );

    // Format context from retrieved chunks
    const formattedContext = this.formatContext(relevantChunks);

    // Generate augmented response using Ollama
    const response = await this.ollamaService.chatWithContext(
      model,
      query,
```

```
        formattedContext,
        options.systemPrompt
    );

    // Add citation metadata to response
    return this.addCitations(response, relevantChunks);
}

// Helper methods for embedding, retrieval, formatting, etc.
}
```

API Endpoints

New API endpoints will be added to support RAG functionality:

Document Management

```
// src/routes/documents.js
const express = require('express');
const router = express.Router();
const multer = require('multer');
const { isAuthenticated } = require('../middleware/auth');

module.exports = function(documentService, ragService) {
    // Configure multer for file uploads
    const upload = multer({
        dest: 'uploads/',
        limits: { fileSize: 50 * 1024 * 1024 } // 50MB limit
    });

    // Upload document
    router.post('/upload', isAuthenticated, upload.single('document'), async (req, res) => {
        try {
            // Create document record
            const document = await documentService.createDocument({
                file: req.file,
                userId: req.user.id,
                collectionId: req.body.collectionId,
                title: req.body.title || req.file.originalname,
                metadata: req.body.metadata ? JSON.parse(req.body.metadata) : {}
            });

            // Save file to Documents directory
            const userDir = path.join('Documents', document.user_id);
            const docDir = path.join(userDir, document.id);

            // Create directories if they don't exist
            if (!fs.existsSync(userDir)) fs.mkdirSync(userDir, { recursive: true });
            if (!fs.existsSync(docDir)) fs.mkdirSync(docDir, { recursive: true });
        }
    });
}
```

```

// Move uploaded file to final destination
const originalExt = path.extname(document.file_name);
const finalPath = path.join(docDir, `original${originalExt}`);
fs.renameSync(req.file.path, finalPath);

// Update document with final path
await documentService.updateDocumentPath(document.id, finalPath);

// Check for duplicates before processing
const isDuplicate = await documentService.checkForDuplicates(document.id);

if (!isDuplicate) {
  // Start processing in background if not a duplicate
  documentService.processDocumentAsync(document.id);
}

res.status(201).json({
  ...document,
  isDuplicate
});
} catch (error) {
  res.status(500).json({ error: error.message });
}
});

// Other document management endpoints
// GET /documents
// GET /documents/:id
// DELETE /documents/:id
// etc.

return router;
};

```

RAG-Enabled Chat

```

// src/routes/chatbot.js (updated)
router.post('/chat-rag', isAuthenticated, async (req, res) => {
  try {
    const { message, sessionId, modelId, useRAG, collectionIds } = req.body;

    // Create chat session if needed
    let session;
    if (sessionId) {
      session = await getSessionById(sessionId, req.user.id);
      if (!session) {
        return res.status(404).json({ error: 'Session not found' });
      }
    } else {
      session = await createNewSession(req.user.id, message.substring(0, 50));
    }
  }
}

```

```
// Save user message
const messageRecord = await saveMessage(req.user.id, message, null,
session.id);

// If RAG is enabled, use RAG service for response
let response;
if (useRAG && collectionIds && collectionIds.length > 0) {
  response = await ragService.queryWithRAG(
    message,
    modelId,
    collectionIds,
    { includeCitations: true }
  );

  // Update message with retrieved chunks
  await updateMessageWithChunks(messageRecord.id, response.citations);
} else {
  // Regular Ollama chat without RAG
  response = await ollamaService.chat(modelId, [{
    role: 'user',
    content: message
  }]);
}

// Update message with AI response
await updateMessageResponse(messageRecord.id, response.text);

res.json({
  sessionId: session.id,
  messageId: messageRecord.id,
  response: response.text,
  citations: response.citations || [],
  model: modelId
});
} catch (error) {
  console.error('Error in chat-rag endpoint:', error);
  res.status(500).json({ error: error.message });
}
});
```

Testing and Validation

The RAG implementation will be tested using:

1. Unit Tests

- Test document processing pipeline
- Test vector store operations
- Test chunking and embedding generation

2. Integration Tests

- Test end-to-end document upload to RAG query
- Test various document types and sizes
- Test vector database connection reliability

3. Performance Testing

- Measure retrieval latency for various collection sizes
- Test embedding generation performance
- Evaluate response quality compared to non-RAG responses

4. User Testing

- Validate relevance of retrieved information
- Test usability of document management interface
- Gather feedback on citation presentation

Embedding Storage and Retrieval

Storage Format

The embeddings will be stored in both the filesystem and ChromaDB:

1. Filesystem Storage

- Location: `Embeddings/user_id/doc_id/`
- Files:
 - `metadata.json`: Document metadata including hash, title, etc.
 - `chunks.json`: Array of text chunks with positions and metadata
 - `embeddings.bin`: Binary file containing embedding vectors

2. ChromaDB Storage

- Collection per user or document group
- Each document chunk stored with:
 - ID: `doc_id:chunk_index`
 - Embedding: Vector representation
 - Metadata: Document info, chunk position, etc.

Embedding Generation

```
// src/services/embeddingService.js
class EmbeddingService {
  constructor(ollamaService, config) {
    this.ollamaService = ollamaService;
    this.embeddingModel = config.embeddingModel || 'nomic-embed-text';
  }

  async generateEmbeddings(texts) {
    if (!Array.isArray(texts)) {
      texts = [texts];
    }
  }
}
```



```
    const embeddings = [];

    // Process in batches to avoid overwhelming the Ollama server
    const batchSize = 10;
    for (let i = 0; i < texts.length; i += batchSize) {
        const batch = texts.slice(i, i + batchSize);
        const batchEmbeddings = await Promise.all(
            batch.map(text =>
this.ollamaService.generateEmbedding(this.embeddingModel, text))
        );
        embeddings.push(...batchEmbeddings);
    }

    return embeddings;
}

async storeEmbeddings(documentId, userId, chunks, embeddings) {
    // Create directory structure
    const userDir = path.join('Embeddings', userId);
    const docDir = path.join(userDir, documentId);

    if (!fs.existsSync(userDir)) fs.mkdirSync(userDir, { recursive: true });
    if (!fs.existsSync(docDir)) fs.mkdirSync(docDir, { recursive: true });

    // Store chunks
    const chunksWithMetadata = chunks.map((chunk, index) => ({
        index,
        content: chunk,
        token_count: this.estimateTokenCount(chunk)
    }));

    fs.writeFileSync(
        path.join(docDir, 'chunks.json'),
        JSON.stringify(chunksWithMetadata, null, 2)
    );

    // Store embeddings in binary format
    const embeddingBuffer = this.convertEmbeddingsToBuffer(embeddings);
    fs.writeFileSync(path.join(docDir, 'embeddings.bin'), embeddingBuffer);

    // Store metadata
    const metadata = {
        document_id: documentId,
        user_id: userId,
        chunk_count: chunks.length,
        embedding_model: this.embeddingModel,
        created_at: new Date().toISOString(),
        dimensions: embeddings[0].length
    };

    fs.writeFileSync(
        path.join(docDir, 'metadata.json'),
        JSON.stringify(metadata, null, 2)
    )
}
```

```
);

return {
  path: docDir,
  metadata
};
}

// Helper methods
convertEmbeddingsToBuffer(embeddings) {
  // Convert array of embedding vectors to binary buffer
  const dimensions = embeddings[0].length;
  const buffer = Buffer.alloc(embeddings.length * dimensions * 4); // 4 bytes
per float32

  embeddings.forEach((embedding, embIndex) => {
    embedding.forEach((value, valueIndex) => {
      buffer.writeFloatLE(value, (embIndex * dimensions + valueIndex) * 4);
    });
  });

  return buffer;
}

loadEmbeddingsFromBuffer(buffer, count, dimensions) {
  // Load embeddings from binary buffer
  const embeddings = [];

  for (let i = 0; i < count; i++) {
    const embedding = [];
    for (let j = 0; j < dimensions; j++) {
      embedding.push(buffer.readFloatLE((i * dimensions + j) * 4));
    }
    embeddings.push(embedding);
  }

  return embeddings;
}

estimateTokenCount(text) {
  // Simple estimation: ~4 characters per token
  return Math.ceil(text.length / 4);
}
}
```

Retrieval Process

When a user query is processed:

1. Query Embedding

```
// Generate embedding for user query
const queryEmbedding = await embeddingService.generateEmbeddings(query);
```

2. Vector Search

```
// Search ChromaDB for similar chunks
const results = await vectorStoreService.search(
  collectionId,
  queryEmbedding,
  {
    limit: topK,
    includeMetadata: true,
    includeEmbeddings: false
  }
);
```

3. Context Assembly

```
// Format retrieved chunks for LLM context
const context = results.map(result => {
  return {
    content: result.metadata.content,
    source: result.metadata.source,
    document_title: result.metadata.document_title,
    document_id: result.metadata.document_id,
    similarity: result.similarity
  };
});

// Format context for LLM prompt
const formattedContext = formatContextForLLM(context);
```

Future Enhancements

1. Advanced Retrieval Methods

- Hybrid search (keyword + vector)
- Re-ranking of retrieved chunks
- Multi-query retrieval strategies

2. Document Processing Enhancements

- OCR for scanned documents and images
- Table extraction and structured data handling
- Automatic metadata extraction

3. Agent Integration

- Combine RAG with Ollama's function calling
- Create specialized agents with domain knowledge
- Implement autonomous research workflows

4. Multi-Model Support

- Use different models for embedding vs. generation
- Support multiple embedding spaces
- Model-specific prompt templates

5. Expanded Knowledge Sources

- Web crawling and automatic updates
- API integrations for external knowledge bases
- Database querying capabilities

6. Embedding Optimization

- Compression techniques for embeddings
- Quantization to reduce storage requirements
- Incremental updates for large documents

Implementation Timeline

| Phase | Timeframe | Key Deliverables |
|-------------------------|-----------|---|
| 1: Foundation | Weeks 1-2 | Database schema, Vector DB integration, Document processing |
| 2: Core RAG | Weeks 3-4 | Query pipeline, Ollama integration, API endpoints |
| 3: Frontend | Weeks 5-6 | Document management UI, Enhanced chat, Knowledge explorer |
| 4: Testing & Refinement | Weeks 7-8 | Testing, Performance optimization, Documentation |